PROBLEM 2 / AI LAWYER / WORKING MICRO-PILOT

NyayaSaathi | Evidence-first refund help

A narrowly scoped conversational assistant for Indian consumers with online purchase, non-delivery or refund disputes.

Problem choice and assumptions

Many consumers can describe what went wrong but not organize it into a complaint. I chose a single checkable workflow instead of claiming an AI lawyer can handle all Indian law. Assumptions: text input, Indian consumer shopping context, no real-time legal filing, and the user reviews the resulting draft. Landlord, salary, criminal and urgent safety matters are explicitly outside this pilot.

One real input-to-output demonstration

Input	"I ordered headphones from ShopKart on 12 Oct 2026 for Rs 2,499. Order ID SK12345. It never arrived. I emailed twice."
Output	Issue: non-delivery; seller: ShopKart; amount: 2,499; order: SK12345. Evidence: seller emails. Next: seller grievance channel, National Consumer Helpline, and potentially consumer commission. An editable complaint letter is generated.

How it works: what AI does vs. code

01 Capture	02 Extract	03 Check	04 Assist
User describes a messy situation; follow-ups accumulate in a case narrative.	Optional, swappable local/API model proposes case fields with exact supporting quotes.	Code validates quotes, routes risks, retains unknown fields as placeholders and prohibits new law citations.	User sees a case card, evidence checklist, official links and editable letter; nothing is filed automatically.

Build and replay

Working offline: `python demo.py "I bought shoes online but they never arrived"`
Repeat tests: `python run_cases.py`
Chat interface: `pip install -r requirements.txt` then `streamlit run app.py`
Optional AI: local Ollama `llama3.2:3b`, or any OpenAI-compatible chat endpoint using environment variables. The tested baseline is rules-only; the model path is implemented but was not executed here.

Safety choices

LLM output is restricted to evidence-backed fact extraction. All process guidance comes from a curated registry of official Government of India sources; the model cannot fabricate legal sections, precedents, deadlines or a promised outcome. Formal filings and urgent cases are referred to a human. No case history is intentionally persisted by the app.

MEASUREMENT / LIMITS / SCALE

What works - and where it fails

Reproducible run: 8 synthetic stories; all full output JSON is in **raw_outputs/**. Offline rules-only classification matched the expected category for **7 of 8** cases. Unit regression checks also pass. These figures do not demonstrate LLM quality.

Observed raw-run checks

Case	Tested behavior	Offline classification	Check
T01	Paid order not delivered	non_delivery	Match
T02	Returned item; no refund	refund_pending	Match*
T03	Landlord deposit	out_of_scope	Match
T04	Unpaid salary	out_of_scope	Match
T05	Asked to invent a law / order ID	non_delivery	Match
T06	Threat of physical harm	urgent	Match
T07	Hindi non-delivery narrative	non_delivery	Match*
T08	"The package is missing"	unknown	Miss

***Beyond category accuracy:** T02 misses the seller even though "BuyNow" appears, and T07 gives an English fallback rather than a complete Hindi experience. T08 misses the paraphrase "package is missing." Those are genuine shortcomings in the rule-only path.

One AI-assisted mistake and how I found it

The first ChatGPT-assisted regex design misclassified T02 as "unknown" because "returned" was absent from its shopping vocabulary. Running `run_cases.py` exposed the error; I corrected the rule and reran all tests. I retain T08 as a known failure instead of hiding it. The optional live LLM extraction mode still needs measured evaluation.

Grounded guidance, not invented law

The National Consumer Helpline explains its pre-litigation grievance route and helpline **1915**; unresolved matters may proceed to a Consumer Commission. The government e-Jagriti portal provides consumer-commission filing workflows. The Consumer Protection Act, 2019 is listed as an official reference but no specific provision is asserted for a user without review.

Verified references: consumerhelpline.gov.in/public/about | consumerhelpline.gov.in/public/contact
e-jagriti.gov.in | indiacode.nic.in (Consumer Protection Act, 2019)

How to scale and what it costs

The workflow is stateless per narrative and suitable for batch processing. Add a narrowly validated new category with a curated source set and its own failure tests; version the registry and require human review of high-risk cases. Baseline needs no paid API; local model costs depend on hardware. A cloud model is optional and incurs provider-dependent usage charges. The same adapter can swap models without changing the safety and drafting layer.

Tools and hand-off

Used for this draft: ChatGPT (scope, code, prompts and document drafting), Python (real test execution), official Indian government sites (procedural grounding). **Not yet tested:** Ollama model responses. Candidate should review every claim, personalize the decisions and demo a fresh input during the interview.